

# CPSC 340:

# Machine Learning and Data Mining

Decision Trees

Bonus slides

# Other Considerations for Food Allergy Example

- What types of preprocessing might we do?
  - Data cleaning: check for and fix missing/unreasonable values.
  - Summary statistics:
    - Can help identify “unclean” data.
    - Correlation might reveal an obvious dependence (“sick”  $\Leftrightarrow$  “peanuts”).
  - Data transformations:
    - Convert everything to same scale? (e.g., grams)
    - Add foods from day before? (maybe “sick” depends on multiple days)
    - Add date? (maybe what makes you “sick” changes over time).
  - Data visualization: look at a scatterplot of each feature and the label.
    - Maybe the visualization will show something weird in the features.
    - Maybe the pattern is really obvious!
- What you do might depend on how much data you have:
  - Very little data:
    - Represent food by common allergic ingredients (lactose, gluten, etc.)?
  - Lots of data:
    - Use more fine-grained features (bread from bakery vs. hamburger bun)?

# Julia Decision Stump Code (not $O(n \log n)$ yet)

Input: feature matrix  $X$  and label vector  $y$

$(n, d) = \text{size}(X)$

$\text{minError} = \text{sum}(y \neq \text{mode}(y))$  compute error if you don't split (user-defined function "mode")

$\text{minRule} = []$

for  $j = 1:d$

for each feature 'j'

for  $i = 1:n$

for each example 'i'

$t = X[i, j]$

set threshold to feature 'j' in example 'i'.

$y_{\text{above}} = \text{mode}(y[X[:, j]. > t])$

find mode of label vector when feature 'j' is above threshold.

$y_{\text{below}} = \text{mode}(y[X[:, j]. \leq t])$

find mode of label vector when feature 'j' is below threshold.

$\hat{y} = \text{fill}(y_{\text{above}}, n)$

Classify all examples based on threshold

$\hat{y}[X[:, j]. \leq t] = y_{\text{below}}$

$\text{error} = \text{sum}(\hat{y} \neq y)$

count the number of errors.

if  $\text{error} < \text{minError}$   
 $\text{minError} = \text{error}$   
 $\text{minRule} = [j \ t]$

store this rule if it has the lowest error so far.

# Going from $O(n^2d)$ to $O(nd \log n)$ for Numerical Features

- Do we have to compute score from scratch?
  - As an example, assume we eat integer number of eggs:
    - So the rules  $(\text{egg} > 1)$  and  $(\text{egg} > 2)$  have same decisions, except when  $(\text{egg} == 2)$ .
- We can actually compute the best rule involving 'egg' in  $O(n \log n)$ :
  - Sort the examples based on 'egg', and use these positions to re-arrange 'y'.
  - Go through the sorted values in order, updating the counts of #sick and #not-sick that both satisfy and don't satisfy the rules.
  - With these counts, it's easy to compute the classification accuracy (see bonus slide).
- Sorting costs  $O(n \log n)$  per feature.
- Total cost of updating counts is  $O(n)$  per feature.
- Total cost is reduced from  $O(n^2d)$  to  $O(nd \log n)$ .
- This is a good runtime:
  - $O(nd)$  is the size of data, same as runtime up to a log factor.
  - We can apply this algorithm to huge datasets.

# How do we fit stumps in $O(nd \log n)$ ?

- Let's say we're trying to find the best rule involving milk:

| Egg | Milk | ... |
|-----|------|-----|
| 0   | 0.7  |     |
| 1   | 0.7  |     |
| 0   | 0    |     |
| 1   | 0.6  |     |
| 1   | 0    |     |
| 2   | 0.6  |     |
| 0   | 1    |     |
| 2   | 0    |     |
| 0   | 0.3  |     |
| 1   | 0.6  |     |
| 2   | 0    |     |

| Sick? |
|-------|
| 1     |
| 1     |
| 0     |
| 1     |
| 0     |
| 1     |
| 1     |
| 1     |
| 0     |
| 0     |
| 1     |

First grab the milk column and sort it (using the sort positions to re-arrange the sick column). This step costs  $O(n \log n)$  due to sorting.

Now, we'll go through the milk values in order, keeping track of #sick and #not sick that are above/below the current value. E.g., #sick above 0.3 is 5.

With these counts, accuracy score is (sum of most common label above and below)/n.

| Milk |
|------|
| 0    |
| 0    |
| 0    |
| 0    |
| 0.3  |
| 0.6  |
| 0.6  |
| 0.6  |
| 0.7  |
| 0.7  |
| 1    |

| Sick? |
|-------|
| 0     |
| 0     |
| 0     |
| 0     |
| 0     |
| 1     |
| 1     |
| 0     |
| 1     |
| 1     |
| 1     |

# How do we fit stumps in $O(nd \log n)$ ?

| Milk | Sick? |
|------|-------|
| 0    | 0     |
| 0    | 0     |
| 0    | 0     |
| 0    | 0     |
| 0.3  | 0     |
| 0.6  | 1     |
| 0.6  | 1     |
| 0.6  | 0     |
| 0.7  | 1     |
| 0.7  | 1     |
| 1    | 1     |

Start with the baseline rule () which is always “satisfied”:

If satisfied, #sick=5 and #not-sick=6.

If not satisfied, #sick=0 and #not-sick=0.

This gives accuracy of  $(6+0)/n = 6/11$ .

Next try the rule (milk > 0), and update the counts based on these 4 rows:

If satisfied, #sick=5 and #not-sick=2.

If not satisfied, #sick=0 and #not-sick=4.

This gives accuracy of  $(5+4)/n = 9/11$ , which is better.

Next try the rule (milk > 0.3), and update the counts based on this 1 row:

If satisfied, #sick=5 and #not-sick=1.

If not satisfied, #sick=0 and #not-sick=5.

This gives accuracy of  $(5+5)/n = 10/11$ , which is better.

(and keep going until you get to the end...)

# How do we fit stumps in $O(nd \log n)$ ?

| Milk | Sick? |
|------|-------|
| 0    | 0     |
| 0    | 0     |
| 0    | 0     |
| 0    | 0     |
| 0.3  | 0     |
| 0.6  | 1     |
| 0.6  | 1     |
| 0.6  | 0     |
| 0.7  | 1     |
| 0.7  | 1     |
| 1    | 1     |

Notice that for each row, updating the counts only costs  $O(1)$ .  
Since there are  $O(n)$  rows, total cost of updating counts is  $O(n)$ .

Instead of 2 labels (sick vs. not-sick), consider the case of 'k' labels:

- Updating the counts still costs  $O(n)$ , since each row has one label.
- But computing the 'max' across the labels costs  $O(k)$ , so cost is  $O(kn)$ .

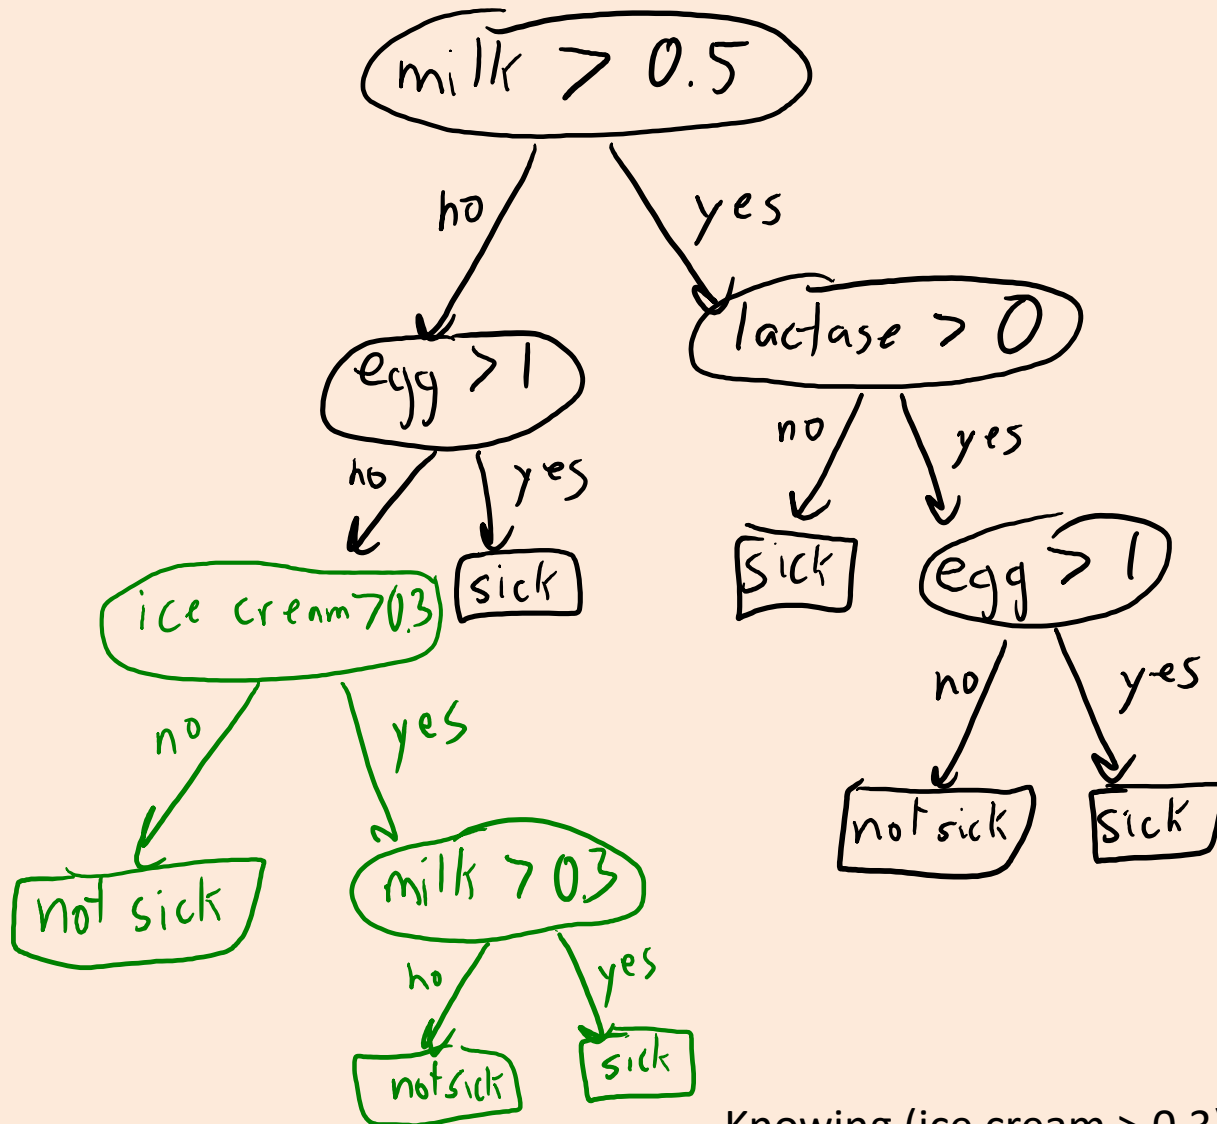
With 'k' labels, you can decrease cost using a "max-heap" data structure:

- Cost of getting max is  $O(1)$ , cost of updating heap for a row is  $O(\log k)$ .
- But  $k \leq n$  (each row has only one label).
- So cost is in  $O(\log n)$  for one row.

Since the above shows we can find best rule in one column in  $O(n \log n)$ ,  
total cost to find best rule across all 'd' columns is  $O(nd \log n)$ .

# Can decision trees re-visit a feature?

- Yes.

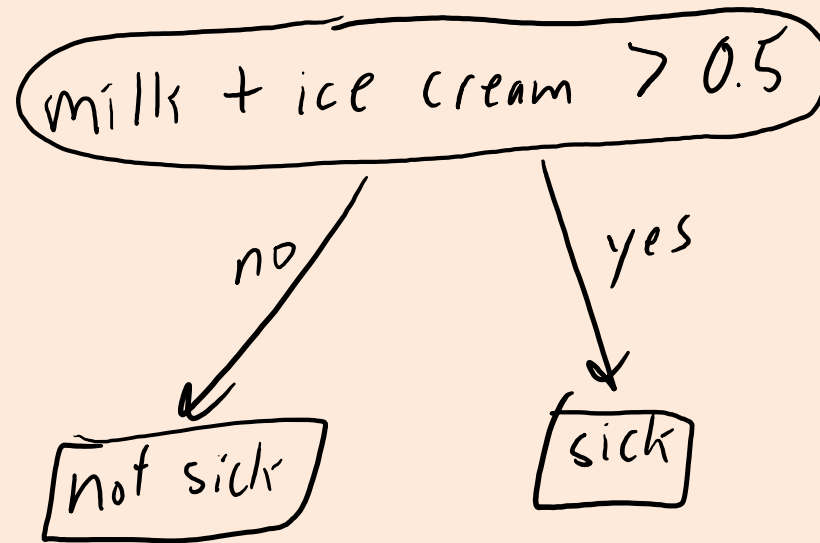


Knowing (ice cream > 0.3) makes small milk quantities relevant.



# Can decision trees have more complicated rules?

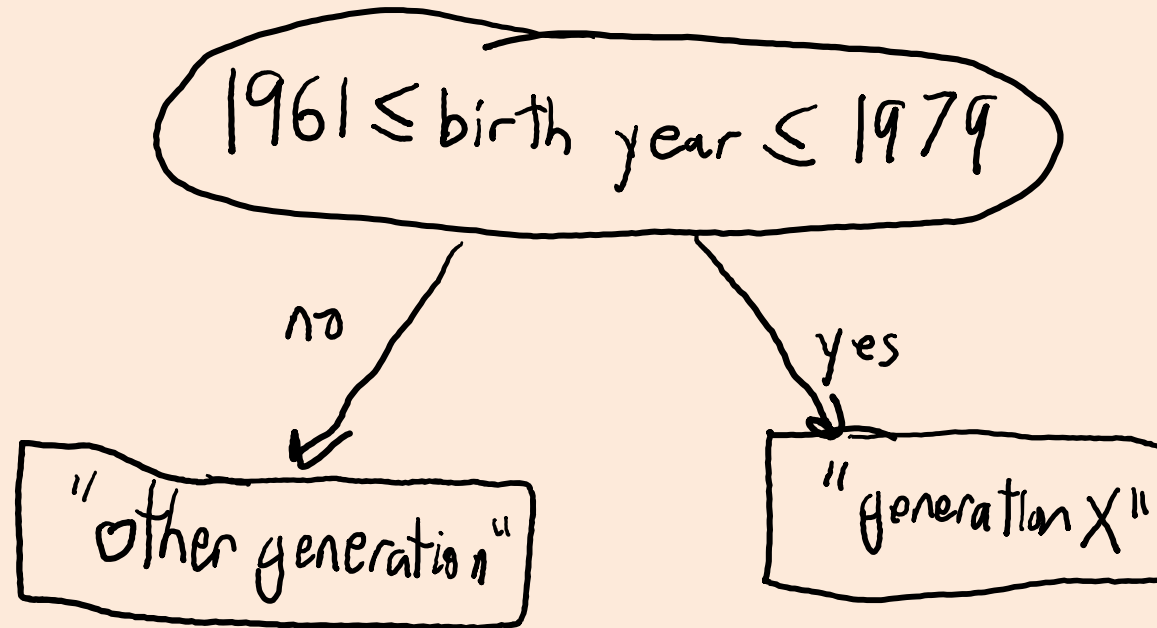
- Yes!
- Rules that depend on **more than one feature**:



- But now searching for the best rule can get expensive.

# Can decision trees have more complicated rules?

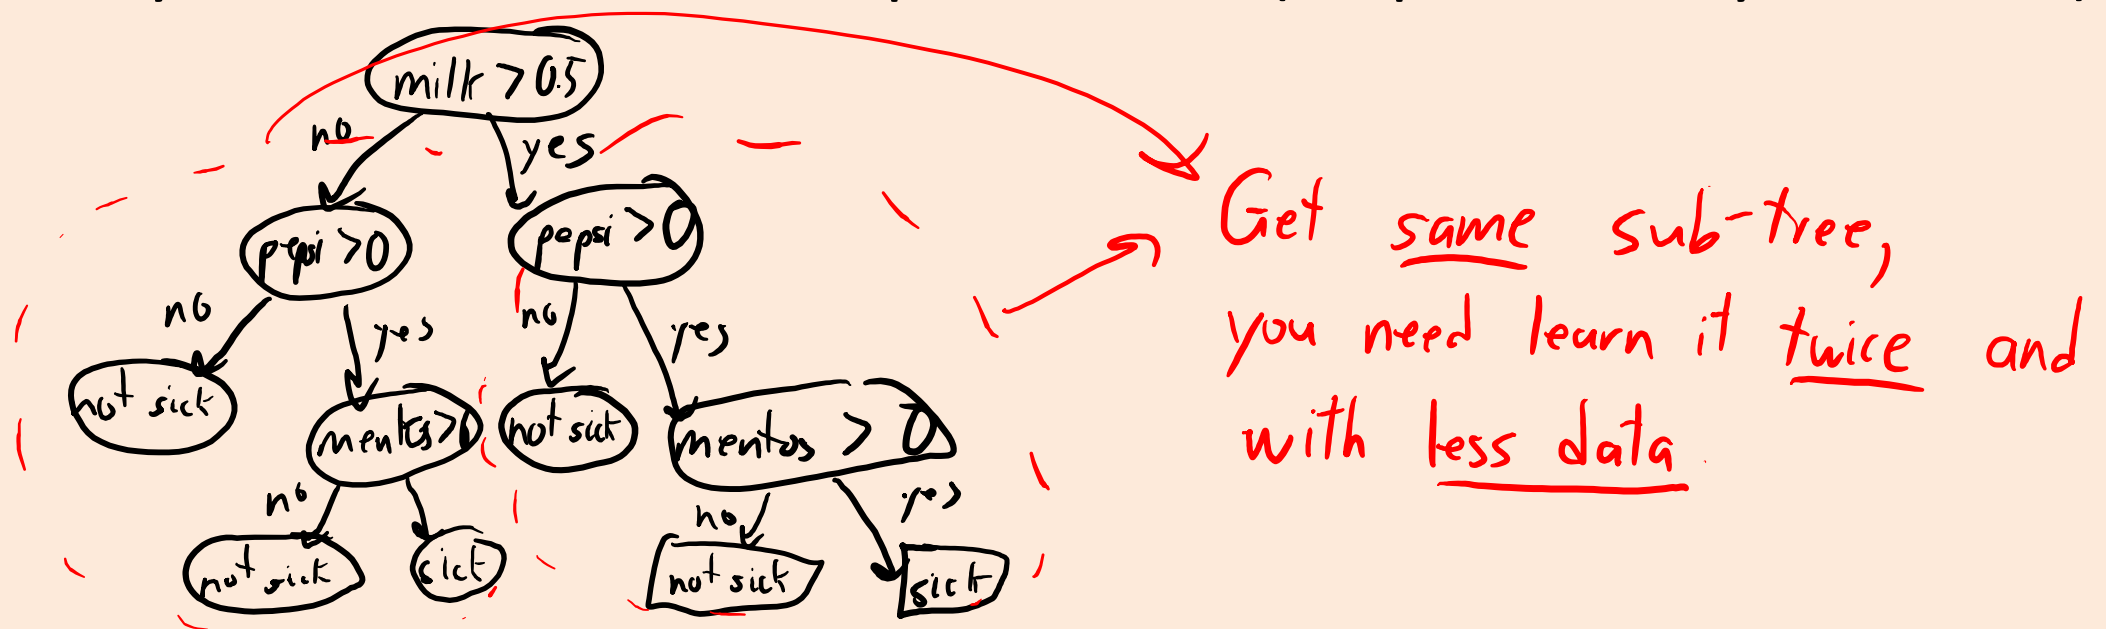
- Yes!
- Rules that depend on **more than one threshold**:



- [“Very Simple Classification Rules Perform Well on Most Commonly Used Datasets”](#)
  - Consider decision stumps based on multiple splits of 1 attribute.
  - Showed that this gives comparable performance to more-fancy methods on many datasets.

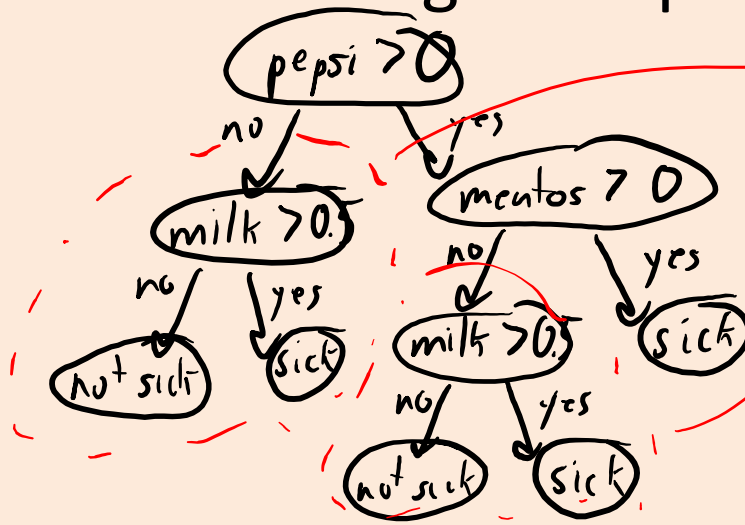
# Does being greedy actually hurt?

- Can't you just go deeper to correct greedy decisions?
  - Yes, but you need to “re-discover” rules with less data.
- Consider that you are allergic to milk (and drink this often), and also get sick when you (rarely) combine diet coke with mentos.
- Greedy method should first split on milk (helps accuracy the most):



# Does being greedy actually hurt?

- Can't you just go deeper to correct greedy decisions?
  - Yes, but you need to “re-discover” rules with less data.
- Consider that you are allergic to milk (and drink this often), and also get sick when you (rarely) combine diet coke with mentos.
- Greedy method should first split on milk (helps accuracy the most).
- Non-greedy method could get simpler tree (split on milk later):



Still has repeated structure,  
but you should <sup>have</sup> more data  
to estimate those splits.

# Decision Trees with Probabilistic Predictions

- Often, we'll have multiple 'y' values at each leaf node.
- In these cases, we might **return probabilities** instead of a label.
- E.g., if in the leaf node we 5 have “sick” examples and 1 “not sick”:
  - Return  $p(y = \text{“sick”} \mid x_i) = 5/6$  and  $p(y = \text{“not sick”} \mid x_i) = 1/6$ .
- In general, a natural estimate of the probabilities at the leaf nodes:
  - Let ' $n_k$ ' be the number of examples that arrive to leaf node 'k'.
  - Let ' $n_{kc}$ ' be the number of times ( $y == c$ ) in the examples at leaf node 'k'.
  - Maximum likelihood estimate for this leaf is  $p(y = c \mid x_i) = n_{kc}/n_k$ .

# Alternative Stopping Rules

- There are more complicated rules for deciding when *\*not\** to split.
- Rules based on **minimum sample size**.
  - Don't split any nodes where the number of examples is less than some 'm'.
  - Don't split any nodes that create children with less than 'm' examples.
    - These types of rules try to make sure that you have enough data to justify decisions.
- Alternately, you can use a **validation set** (see next lecture):
  - Don't split the node if it decreases an approximation of test accuracy.